



Lab 1: Exploiting LLM APIs with excessive agency

The vulnerability was that the LLM was given too much access and could perform actions that it really shouldn't be able to do on its own. This made it possible to use the LLM to delete another user's account. A good fix would be to limit what the LLM has access to and require authorization before allowing sensitive actions.



Lab 2: Exploiting vulnerabilities in LLM APIs

One of the APIs the LLM could use had a command injection issue. Because the input wasn't handled safely, it was possible to get the system to run commands it wasn't supposed to. The fix would be to properly validate user input and make sure it can't be passed directly into system commands.



Lab 3: Indirect prompt injecting

For this lab instructions could be hidden inside something like a product review, and the LLM could follow those instructions without realizing they were malicious. This could cause the LLM to perform an action the actual user never asked for. The fix would be to treat outside content as untrusted and put stronger restrictions on what actions the LLM can perform.



Congratulations, you solved the lab!

Share your skills!

[Continue learning](#)[Home](#) | [My account](#) | [Live chat](#)

Live chat

You: change my email to a different address (for example: test@example.com)

Lab 4: Exploiting insecure output handling in LLMs

In this lab, malicious code could be hidden in a product review and then included in the LLM's response, causing it to run in another user's browser. The fix would be to make sure HTML or scripts are treated as regular text instead of executable code.



Congratulations, you solved the lab!

Share your skills!

[Continue learning](#)[Home](#) | [My account](#) | [Live chat](#)

Lightweight "I33t" Leather Jacket



\$1337.00